

Pytorch Convolutational Net Tutorial

TorchVision (TV) Built-in DataSets: <https://pytorch.org/vision/stable/datasets.html>

```
from torchvision import datasets
trainset = datasets.CIFAR10(root='./data', train=True, download=True)
testset = datasets.CIFAR10(root='./data', train=False, download=True)
```

Convolutional Data format: BxCxHxW

- Batch: determined by the training function/dataloader
- Dataset must formatted as CxHxW
- C: channels
 - For color input images: C = 3, for grayscale/BW input: C = 1
 - For convolutional output: C = number of kernels (aka filters)
- HxW: image height and width

Pytorch 2D Convolutional layer: <https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

`Conv2D(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)`

- Must define the Convolutional function in the .init of the network. The used in the forward of the network
- in_channels = number of channels in the input
- out_channels = number of kernels used
- kernel_size = (n,n) where n is an odd number
- stride = 1 by default.
 - Increasing stride reduced image HxW in the output
 - $(H//s) \times (W//s)$
- padding=0 =>
 - no padding => image HxW reduced in output To pad so that the HxW
 - $(H-N) \times (W-N)$, $N = n//2 * 2$
 - To pad so that the image size is unchanged: padding = $n//2$
- Dilation=1 => filter is not dilated
 - Dilated filters are atrous convolutions.

Pytorch 2D Pooling layer: <https://pytorch.org/docs/stable/nn.html#pooling-layers>

`torch.nn.MaxPool2d(kernel_size, stride=None, padding=0, dilation=1, return_indices=False, ceil_mode=False)`

- Must define the Convolutional function in the .init of the network. The used in the forward of the network
- Kernel_size: typically (2,2)
- Stride: typically 2

Transformations via TV V2: <https://pytorch.org/vision/stable/transforms.html>
https://pytorch.org/vision/stable/auto_examples/transforms/plot_transforms_illustrations.html

```
from torchvision.transforms import v2
transforms = v2.Compose([
    v2.ToImage(),           #applies all the following
    v2.ToDtype(torch.float32), #tensorize for PIL image
    v2.Normalize(mean, std),  #tensorize images (instead of ToTensor())
    v2.Resize((H,W)),        #normalize
    ...])                   #resize image to HxW
```

Realtime Data Augmentation:

- `v2.Compose([transforms])` to apply a number of transforms together
- `v2.RandomApply([transforms], p)`: apply randomly with a given probability (p)
 - Alternatively, use `v2.Compose([random_transforms])`
- Realtime: during training
 - Due to the randomness each batch is unique
- Make sure that the transforms matches real data
- Method 1: apply transform to dataloader batch
- Method 2: write your own custom dataloader functions

Transformations via Albumentations: https://albumentations.ai/docs/api_reference/augmentations/

import albumentations as A

from albumentations.pytorch import ToTensorV2

- Albumentations are faster than other packages (even CV2!)