



Mid-Level CV: Feature Detection

ENEE 4584/5584 CV app in DL

Dr. Alsamman

Slide Credits:



Image Features

- ❖ Information in an image that is unique to the image
- ❖ Global features
 - One feature vector per image
 - Examples: histograms, shape, texture, eigenspace, etc.
 - Allow for compact representation of an image
 - Standard algorithms can be used
 - Sensitive to occlusion, clutter, deformations, etc.
 - Segmentation must be used.



Image Features

❖ Local features

- Feature as a patch or region of interest.
- No need for segmentation.
- Robust to occlusion, clutter, deformations, etc.
- Multiple feature vectors of varying sizes.
- Non-standard algorithms



Local Features Desired Properties

❖ Repeatability

- The same feature can be found in several images despite geometric and photometric transformations

❖ Saliency

- Each feature has a distinctive description

❖ Compactness and efficiency

- Many fewer features than image pixels

❖ Locality

- A feature occupies a relatively small area of the image;
- robust to clutter and occlusion



Popular Features

❖ Harris Corner Detection [1989]

➤ https://docs.opencv.org/3.4/dc/d0d/tutorial_py_features_harris.html

➤ Shi-Tomasi Corner Detector [1994]

▪ https://docs.opencv.org/3.4/d4/d8c/tutorial_py_shi_tomasi.html

❖ Scale-Invariant Feature Transform (SIFT) [2004]

➤ https://docs.opencv.org/3.4/da/df5/tutorial_py_sift_intro.html

➤ Speeded-Up Robust Features (SURF) [2006]

▪ https://docs.opencv.org/3.4/df/dd2/tutorial_py_surf_intro.html

❖ Features from Accelerated Segment Test (FAST) [2006,2010]

➤ https://docs.opencv.org/3.4/df/d0c/tutorial_py_fast.html

❖ Binary Robust Independent Elementary Features (BRIEF) [2010]

➤ https://docs.opencv.org/3.4/dc/d7d/tutorial_py_brief.html

❖ Oriented FAST and Rotated BRIEF (ORB) [2011]

➤ https://docs.opencv.org/3.4/d1/d89/tutorial_py_orb.html





Harris Corner Detection



Corner Detection

- ❖ Local features that are insensitive to deformation (scale, rotation).
 - Use for detection, calibration, etc.
- ❖ Edge detection is not very good for corner detection.
 - Focuses on gradient in x-y direction.
 - Corner is a discontinuity: gradient is not defined.
- ❖ Idea: shifting a window in the image in multiple directions
 - Max difference when there is a corner





Windowed Square Sum Difference (SSD)

❖ Defined an a patch/window size

- Size = $2m \times 2m$
- Can use a Gaussian window

❖ Windowed SSD

- Shift the window by u, v

$$E(x, y) = \sum_{m, n} w(m, n) [I(x + u + m, y + v + n) - I(x, y)]^2$$

↑
change

↖
weighting window

↑
shifted image

↗
original image

$$E(x, y) = w(x, y) * [I(x + u, y + v) - I(x, y)]^2$$

↗
2D convolution



Taylor Series Approximation

$$f(x+u, y+v) = f(x, y) + uf_x(x, y) + vf_y(x, y) +$$

First partial derivatives

$$\frac{1}{2!} [u^2 f_{xx}(x, y) + uv f_{xy}(x, y) + v^2 f_{yy}(x, y)] +$$

Second partial derivatives

$$\frac{1}{3!} [u^3 f_{xxx}(x, y) + u^2 v f_{xxy}(x, y) + uv^2 f_{xyy}(x, y) + v^3 f_{yyy}(x, y)]$$

Third partial derivatives

+ ... (Higher order terms)

$$f(x+u, y+v) \approx f(x, y) + uf_x(x, y) + vf_y(x, y)$$



Taylor series expansion

For small shifts

$$\begin{array}{l} u \rightarrow 0 \\ v \rightarrow 0 \end{array} \Rightarrow I(x+u, y+v) \approx I(x, y) + \frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v$$

$$I(x+u, y+v) \approx I(x, y) + [I_x \quad I_y] \begin{bmatrix} u \\ v \end{bmatrix}$$

image derivatives
along x and y axes



SSD Derivation

$$E(x, y) = w(x, y) * \left(I(x, y) + [I_x \quad I_y] \begin{bmatrix} u \\ v \end{bmatrix} - I(x, y) \right)^2$$

$$= w(x, y) * \left([I_x \quad I_y] \begin{bmatrix} u \\ v \end{bmatrix} \right)^2$$

$$= w(x, y) * \left([I_x \quad I_y] \begin{bmatrix} u \\ v \end{bmatrix} \right)^T \left([I_x \quad I_y] \begin{bmatrix} u \\ v \end{bmatrix} \right)$$



$$E(x, y) = w(x, y) * [u \quad v] \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \quad I_y] \begin{bmatrix} u \\ v \end{bmatrix}$$

$$= [u \quad v] \underbrace{\left(w(x, y) * \begin{bmatrix} I_x^2(x, y) & I_x(x, y)I_y(x, y) \\ I_x(x, y)I_y(x, y) & I_y^2(x, y) \end{bmatrix} \right)}_{M(x, y)} \begin{bmatrix} u \\ v \end{bmatrix}$$



$$E(x, y) = [u \quad v] M(x, y) \begin{bmatrix} u \\ v \end{bmatrix}$$

For a Gaussian window

↙

$$M(x, y; \sigma) = w(x, y; \sigma) * \begin{bmatrix} I_x^2(x, y) & I_x(x, y)I_y(x, y) \\ I_x(x, y)I_y(x, y) & I_y^2(x, y) \end{bmatrix}$$

$$M(x, y; \sigma) = \begin{bmatrix} w(x, y; \sigma) * I_x^2(x, y) & w(x, y; \sigma) * I_x(x, y)I_y(x, y) \\ w(x, y; \sigma) * I_x(x, y)I_y(x, y) & w(x, y; \sigma) * I_y^2(x, y) \end{bmatrix}$$



Image Gradient


$$I_x(x, y) = I(x + 1, y) - I(x, y)$$

$$= \underbrace{\begin{bmatrix} 0 & 0 & 0 \\ 0 & -1 & 1 \\ 0 & 0 & 0 \end{bmatrix}}_{D_x(x,y)} * I(x, y)$$




$$I_y(x, y) = I(x, y + 1) - I(x, y)$$


$$= \underbrace{\begin{bmatrix} 0 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 1 & 0 \end{bmatrix}}_{D_y(x,y)} * I(x, y)$$






I_x



I_y




$I_x I_y$





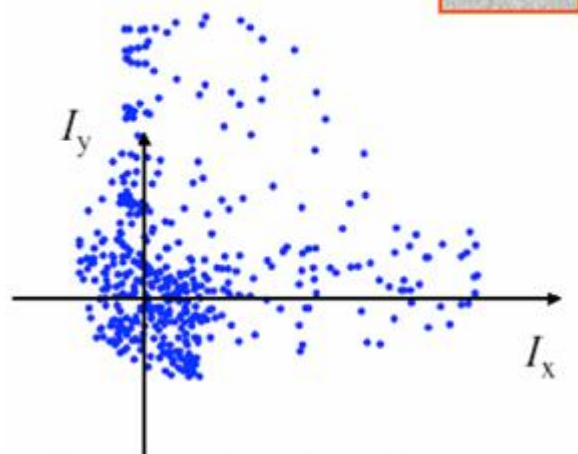

Matrix Form

- ❖ Automated method for edge detection
- ❖ 2nd order Taylor series approx of SSD:
 - Matrix form: $S \approx \frac{1}{2} [u \ v] M [u \ ; \ v]$
 - $M = \sum_x \sum_y w(x,y) * [I_x^2 \ I_x I_y \ ; \ I_x I_y \ I_y^2]$
 - I_x, I_y = gradient in x, y
- ❖ $S < \text{Threshold}$ is an ellipse
 - Ellipse can be described by the orthogonal basis of M
 - Orthogonal basis: direction of fastest and slowest change
 - Use eigenvectors of M to detect the 2 orthogonal basis: v_1, v_2
 - $\lambda_1 > \lambda_2 \Rightarrow v_1$ is direction of fastest change






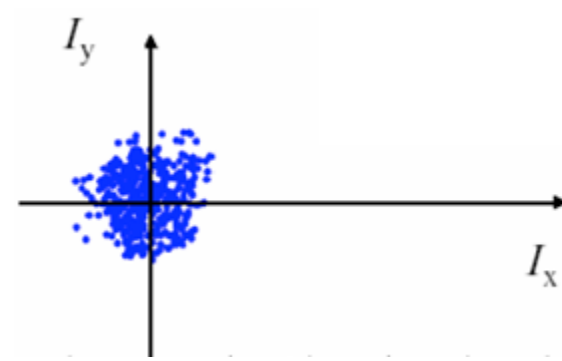
Corner



Linear Edge



Flat




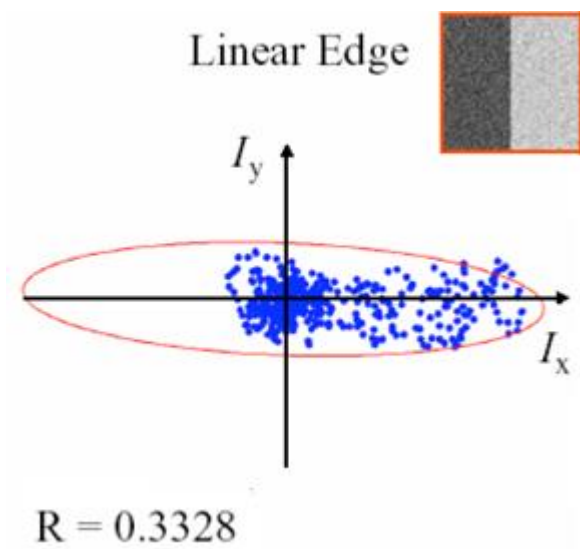
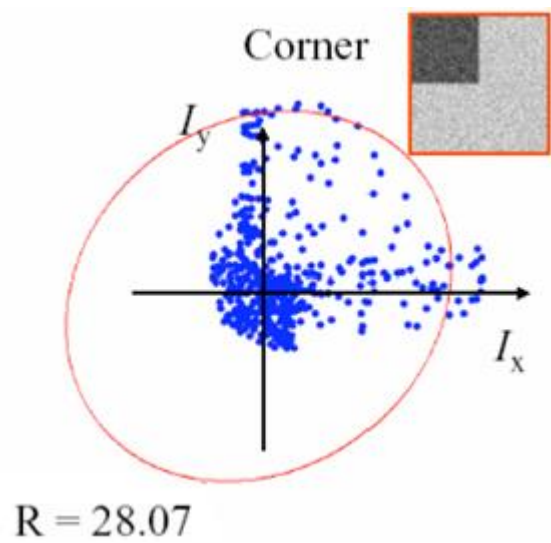


❖ Corners, edges, flat regions:

- Edge point: $\lambda_1 > \lambda_2$ or $\lambda_2 > \lambda_1$
- Flat regions: small λ_1, λ_2
- Corners: $\lambda_1 \approx \lambda_2$; λ_1, λ_2 large

❖ λ 's can be used to generate a response parameter, R:

- $R = \lambda_1 \lambda_2 - k(\lambda_1 - \lambda_2)^2$
- $k = \text{constant } (\sim 0.04-0.06)$
- Edges: $R < 0$
- Flat regions: $|R|$ is small
- Corners: $R > 0$





Harris Corner Detection Algorithm

- ❖ Smooth image to reduce noise
 - Harris is susceptible to noise.
- ❖ Compute magnitude of x, y gradient at each pixel
 - Remember images origin is upper-left corner.
 - Can be performed in a window. Harris uses Gaussian window.
- ❖ Find eigen values of A
 - Can be computed in closed form: $A = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$
$$\lambda_1 = (a+d - \sqrt{(a+d)^2 + 4bc})/2$$
$$\lambda_2 = (a+d + \sqrt{(a+d)^2 + 4bc})/2$$
- ❖ Check the λ_1, λ_2 .
 - A response function can be used: $R = \lambda_1\lambda_2 - k(\lambda_1 - \lambda_2)^2$
 - $k = \text{constant } (\sim 0.04-0.06)$
- ❖ Keep local maxima of R







SIFT



SIFT Features

❖ Scale Invariant Feature Transform

- Detects multiple blob sizes at various scales
- Keeps blobs with the greatest response at each scale
- Blob center located with sub-pixel accuracy

❖ Extracts local feature descriptors



- Blobs are converted into a feature
- Within each blob a gradients and their orientation are compiled into a histograms
- Histograms provides some invariance to rotation
- A feature vector descriptor makes matching a features easy



❖ *Reasonably* invariant to changes in illumination, image noise, rotation, scaling, and small changes in viewpoint.






LoG Blob Detection

- ❖ Blobs have spatial size (scale)
 - Laplacian looks for points (independent of scale)
- ❖ Blobs tend to have a gaussian spread.
- ❖ By convolving the Laplacian with a Gaussian
 - AKA Laplacian of Gaussian (LoG)
 - Gaussian will blur smaller blobs/points
 - Laplacian will highlight edges
- ❖ Response is maximized when $\sigma = r/\sqrt{2}$
 - r is radius of the blob









Laplacian of Gaussian (LoG)

$$\nabla^2 g(x, y) = \frac{\partial^2 g(x, y)}{\partial x^2} + \frac{\partial^2 g(x, y)}{\partial y^2}$$



$$\nabla^2 g(x, y) = \left(\frac{x^2}{\sigma^4} + \frac{y^2}{\sigma^4} - \frac{2}{\sigma^2} \right) \exp \left(-\frac{x^2 + y^2}{2\sigma^2} \right)$$

❖ The response is max when $\nabla^2 G=0$ or:

$$\left(\frac{x^2}{\sigma^4} + \frac{y^2}{\sigma^4} - \frac{2}{\sigma^2} \right) = 0 \Rightarrow \frac{x^2 + y^2}{\sigma^4} = \frac{2}{\sigma^2}$$

$$x^2 + y^2 = 2\sigma^2 \Rightarrow r^2 = 2\sigma^2$$

$$\sigma = \frac{r}{\sqrt{2}}$$





Difference of Gaussian

- ❖ LoG can be approximated by a DoG
- ❖ Subtract two Gaussians at slightly different scales
 - Correlates well with our understanding of how human vision works


$$D(x, y) = \frac{1}{2\pi\sigma_1^2} \exp\left(-\frac{x^2 + y^2}{2\sigma_1^2}\right) - \frac{1}{2\pi\sigma_2^2} \exp\left(-\frac{x^2 + y^2}{2\sigma_2^2}\right)$$

- $\sigma_1 > \sigma_2$.


- $\sigma_1 = k\sigma_2$




Difference of Gaussians





-





=








scale →					
octave	0.707107	1.000000	1.414214	2.000000	2.828427
	1.414214	2.000000	2.828427	4.000000	5.656854
	2.828427	4.000000	5.656854	8.000000	11.313708
	5.656854	8.000000	11.313708	16.000000	22.627417






Scale-space Extrema

- ❖ Local maxima/minima are detected in 3x3x3 neighborhoods
- ❖ Spans adjacent DoG images across scales.
- ❖ Interpolation of nearby data determine its position.
- ❖ Keypoints with low contrast
- ❖ Responses along edges are
- ❖ The keypoint is assigned an
 - gradient orientation histogram









Reject Low Contrast Points

❖ Reject points with bad contrast:

➤ $D < 0.03$ (image values $[0,1]$)





Reject Points along Edges

- ❖ Problem: Strong responses along edges
- ❖ Solution eliminate edge points
- ❖ Use Hessian to determine edges

$$H = \begin{bmatrix} I_x^2 & I_{xy} \\ I_{xy} & I_y^2 \end{bmatrix}$$

- $I_x = D \otimes \begin{bmatrix} -1 \\ 1 \end{bmatrix}$ = derivative along x
- $I_y = D \otimes \begin{bmatrix} -1 & 1 \end{bmatrix}$ = derivative along y
- $I_x^2 = I_x I_x$,
- $I_y^2 = I_y I_y$,
- $I_{xy} = I_x I_y$





❖ Edge point: $\lambda_1 \gg \lambda_2$

➤ Use a measure :

$$r = \frac{\text{trace}^2}{\text{determinant}} = \frac{(\lambda_1 + \lambda_2)^2}{\lambda_1 \lambda_2}$$

➤ reject point if : $\lambda_1 > 10\lambda_2$:

$$r > \frac{(10 + 1)^2}{10} = 12.1$$





SURF




Speed Up Robust Feature (SURF)


- ❖ Improvements to SIFT
 - SIFT used DoG
 - 9x9 Haar-like wavelet response
 - Convolution with box filter can be easily calculated from integral images.
- ❖ Scale up the filter instead of resizing image
 - Double filter
 - Search for extrema in 3x3x3 neighborhood.
- ❖ Determinant of Hessian matrix for both scale and location
 - $\det\{H\} = I_{xx}I_{yy} - (wI_{xy})^2$
 - $w = 0.9$




$I_x I_x$



$I_y I_y$



$I_x I_y$






FAST



Features from accelerated Segment Test (FAST)

- ❖ Focuses on real-time application
- ❖ Corner detection technique
- ❖ Criteria for corner:
 - Each pixel, p_i , has a circle of 16 pixels around it, p_{i_1} to $p_{i_{16}}$
 - Radius = 3
 - p_i is a corner if at least 12, p_{i_j} are brighter or darker than a threshold

$$|p_i - p_{i_j}| > Threshold$$
- ❖ High speed test:
 - Examine only the four pixels at $j = 1, 9, 5$ and 13
 - At least 3 must be brighter/darker





Machine Learning

- ❖ Machine learning for better performance
- ❖ Weaknesses of previous algorithm: High acceptance rate: not inclusive of all j candidates
- ❖ Solution: Train for optimal performance
 1. Select a training data set
 - Identify corners
 2. Run FAST for all 16 neighbors on a circle
 - Store all 16 neighbor values
 3. Process neighbors into 3 states: darker ($d=-1$), similar (0), brighter (1)
 4. Use decision tree classifier
 - find the neighbors that are most important to determine if p_i is a corner
 5. Use decision tree for test images



BRIEF



Descriptor Vectors

- ❖ Feature vectors can be extremely large
 - SIFT: 128 values to describe each feature
 - Values are floating point numbers
 - Single precision = 4B/value, Double precision = 8 B/value
- ❖ Use dimension reduction techniques
 - Principal component analysis (PCA)
 - Linear discriminant analysis (LDA)
- ❖ Use hashing methods to convert floating point values into binary strings
 - Locality Sensitive Hashing (LSH)
 - Allows for fast comparison of features: XOR then count bits



Binary Robust Independent Elementary Features (BRIEF)

- ❖ Detecting a feature
- ❖ Create an $S \times S$ patch around the feature
- ❖ Randomly choose n random pairs in $S \times S$
 - Uniform random selection works well
 - Gaussian works well
 - Polar
- ❖ Each pair (i,j) is given a binary value
 - $P(i) \geq P(j) \Rightarrow 0$
 - $P(i) < P(j) \Rightarrow 1$
- ❖ Scale each by 2^k



ORB



Oriented FAST and Rotated BRIEF

- ❖ Developed by OpenCV Labs
 - No patented as SIFT & SURF are patented
- ❖ Use FAST to find corners
 - Apply to many scales as in SIFT
- ❖ Use Harris score to pick top corners
- ❖ New: orientation of corner
 - From corner center to weighted patch center
 - Discretized to $2\pi/30$ (12)
- ❖ Use BRIEF for feature vectors