

09: GLIBC HEAP EXPLOITS

Vassil Roussev

vassil@cs.uno.edu

code/4621-heap.zip


REF

- <https://github.com/shellphish/how2heap>
- <https://azeria-labs.com/heap-exploitation-part-2-glibc-heap-free-bins/>
- <https://heap-exploitation.dhavalkapil.com/>

GLIBC HEAP OPERATION

THE HEAP

- Dynamically allocated/deallocated memory
 - » *on request*
- Allocations managed by **glibc** (Linux)
 - » *fast performance + efficiency is critical*
 - » *memory allocator maintains metadata*
 - » *overall, a non-trivial management algorithm*
- API
 - » **malloc** / **calloc** (realloc / reallocarray)
 - » **free**

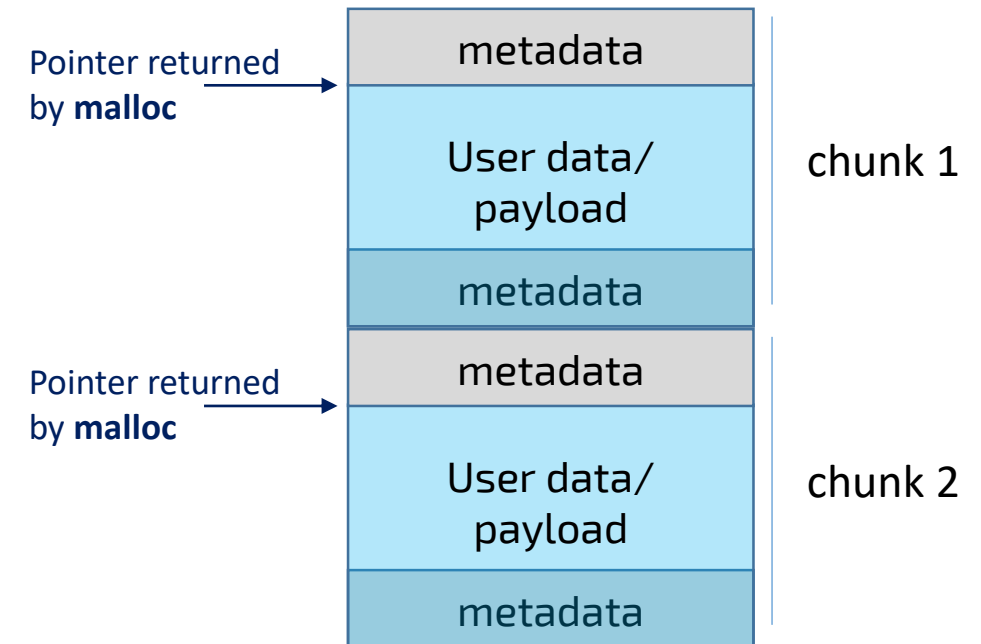


MEMORY CHUNKS

```
struct malloc_chunk {
    INTERNAL_SIZE_T      mchunk_prev_size; /* Size of previous chunk, if it is free. */
    INTERNAL_SIZE_T      mchunk_size;      /* Size in bytes, including overhead. */
    struct malloc_chunk* fd;                /* double links -- used only if this chunk is free. */
    struct malloc_chunk* bk;
    /* Only used for large blocks: pointer to next larger size. */
    struct malloc_chunk* fd_nextsize; /* double links -- used only if this chunk is free. */
    struct malloc_chunk* bk_nextsize;
};

typedef struct malloc_chunk* mchunkptr;
```

- Allocated chunks are
 - » *preceded by metadata header*
 - » *aligned on 8-byte (32-bit) / 16-byte (64-bit) boundaries*
- Free'd chunks are organized into **bins** (linked lists)
 - » 10x **FAST BINS** (fixed size): 16, 24, ..., 88 bytes + header
 - » 1x **UNSORTED BIN**
 - » 62x **SMALL BINS**: 16, 24, ...504
 - neighbors may be coalesced & placed in the **unsorted**
 - » 63x **LARGE BINS**
 - 32x size 64
 - 16x size 512
 - 8x size 4096
 - 4x size 32768
 - 2x size 262144
 - 1x everything else



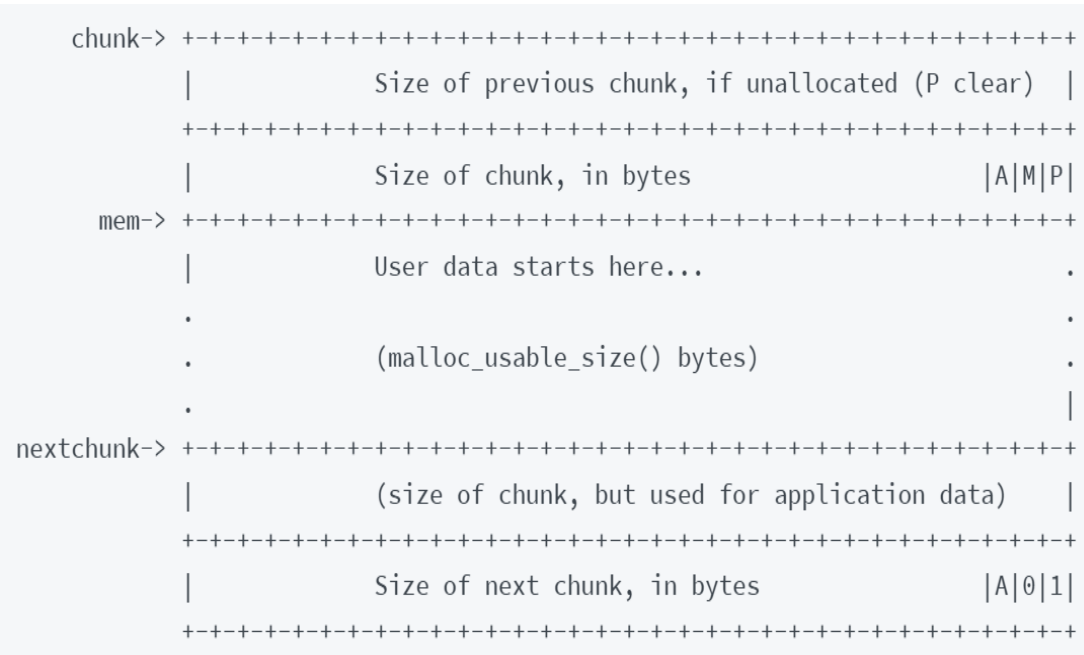
CHUNK LAYOUT

```

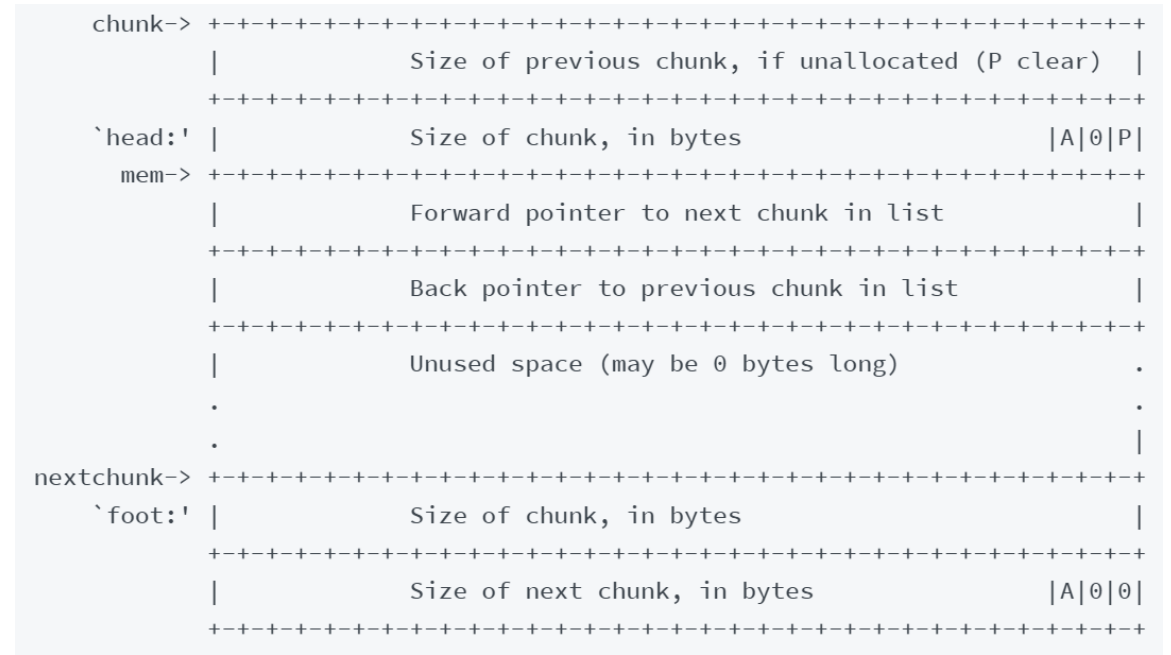
struct malloc_chunk {
    INTERNAL_SIZE_T      mchunk_prev_size; /* Size of previous chunk, if it is free. */
    INTERNAL_SIZE_T      mchunk_size;      /* Size in bytes, including overhead. */
    struct malloc_chunk* fd;                /* double links -- used only if this chunk is free. */
    struct malloc_chunk* bk;
    /* Only used for large blocks: pointer to next larger size. */
    struct malloc_chunk* fd_nextsize; /* double links -- used only if this chunk is free. */
    struct malloc_chunk* bk_nextsize;
};

typedef struct malloc_chunk* mchunkptr;

```



allocated



free

HEAP RULES / VIOLATIONS

- **Do not** read/write to a **malloc**-ed pointer after it has been **free**-d
 - » *Use after free vulnerability*
- **Do not** use or leak uninitialized info in a heap allocation
 - » *Information leaks & uninitialized data vulnerabilities*
- **Do not** read/write bytes after the end of an allocation
 - » *Heap overflow & read beyond bounds vulnerabilities*
- **Do not** pass a pointer that originated from **malloc** to **free** more than once
 - » *Double free vulnerability*
- **Do not** read/write bytes before the beginning of an allocation
 - » *Heap underflow vulnerability*
- **Do not** pass a point that did not originate from **malloc** to **free**
 - » *Invalid free vulnerability*
- **Do not** use a pointer returned by **malloc** before checking against **NULL**
 - » *Null-dereference bug & occasional arbitrary write vulnerability*

HOW DOES **FREE** WORK? (SIMPLE VERSION)

- `free(NULL)` → a valid no-op
- Sanity checks
 - » *address alignment*
 - » *is size realistic (not huge, not too small)*
 - » *is address within **arena** boundaries*
 - » *is chunk not already free*
 - » ...
- Bottom line
 - » *these are sanity check that an attacker can readily bypass*

FREED CHUNKS

- Freed chunks are
 - » *added to free lists*
 - » *used to store metadata about the lists*
 - like **boundary tags**

```
struct malloc_chunk {  
    INTERNAL_SIZE_T      mchunk_prev_size; /* Size of previous chunk, if it is free. */  
    INTERNAL_SIZE_T      mchunk_size;      /* Size in bytes, including overhead. */  
    struct malloc_chunk* fd;                /* double links -- used only if this chunk is free. */  
    struct malloc_chunk* bk;  
    /* Only used for large blocks: pointer to next larger size. */  
    struct malloc_chunk* fd_nextsize; /* double links -- used only if this chunk is free. */  
    struct malloc_chunk* bk_nextsize;  
};  
  
typedef struct malloc_chunk* mchunkptr;
```

A (0x04) Allocated arena
M (0x02) Mmap'd chunk
P (0x01) Prev chunk in use

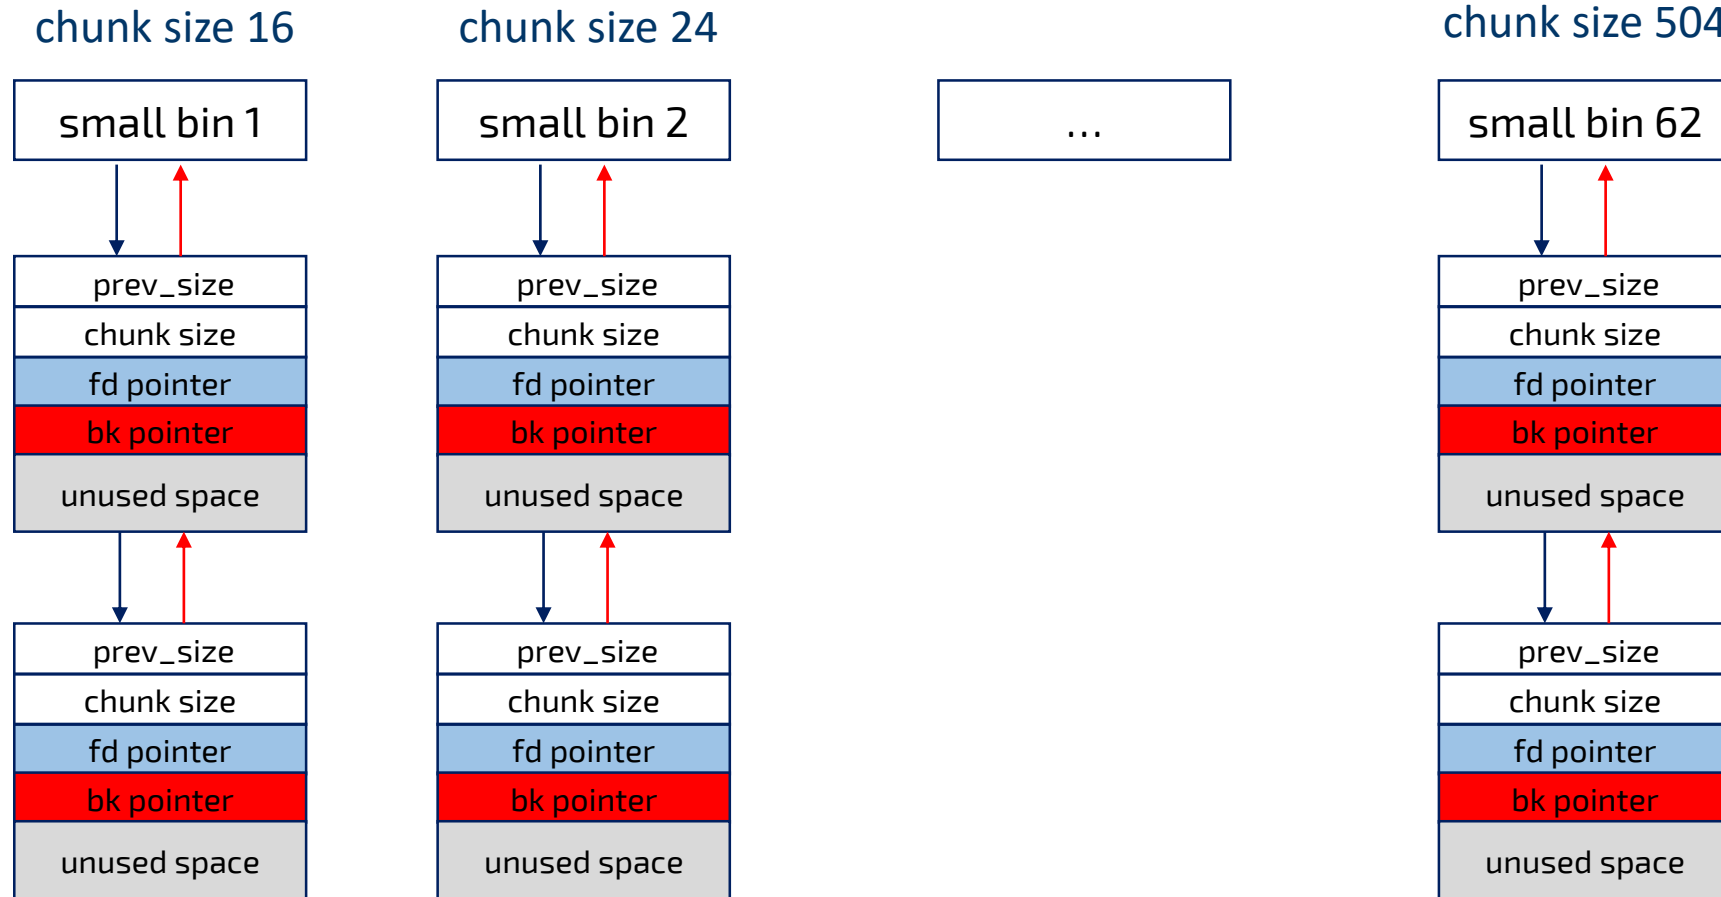
chunk 1

chunk size	A	M	P
fwd pointer			
bck pointer			
fd_nextsize			
bk_nextsize			
...			
prev_size			

RECYCLING MEMORY W/ BINS

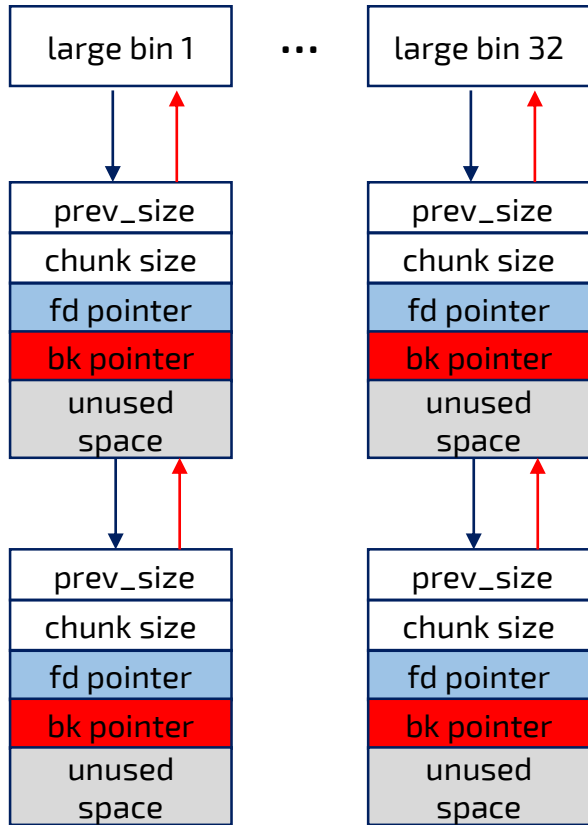
- Bins == lists
 - » BIN[0] → N/A
 - » BIN[1] → UNSORTED
 - » BIN[2] – BIN[64] → SMALL BINS
 - » BIN[65] – BIN[127] → LARGE BINS
- Baseline **free** algorithm
 - » *If the chunk has the M bit set*
 - allocation was allocated off-heap and should be **munmapped**.
 - » *Otherwise*
 - if the chunk before this one is free, the chunk is merged backwards to create a bigger free chunk.
 - » *Similarly*
 - if the chunk after this one is free, the chunk is merged forwards to create a bigger free chunk.
 - » *If this potentially-larger chunk borders the “top” of the heap*
 - the whole chunk is absorbed into the end of the heap, rather than stored in a “bin”.
 - » *Otherwise*
 - the chunk is marked as free and placed in an appropriate bin.

SMALL BINS (< 512 / 1024 BYTES ON 32-/64-BIT)

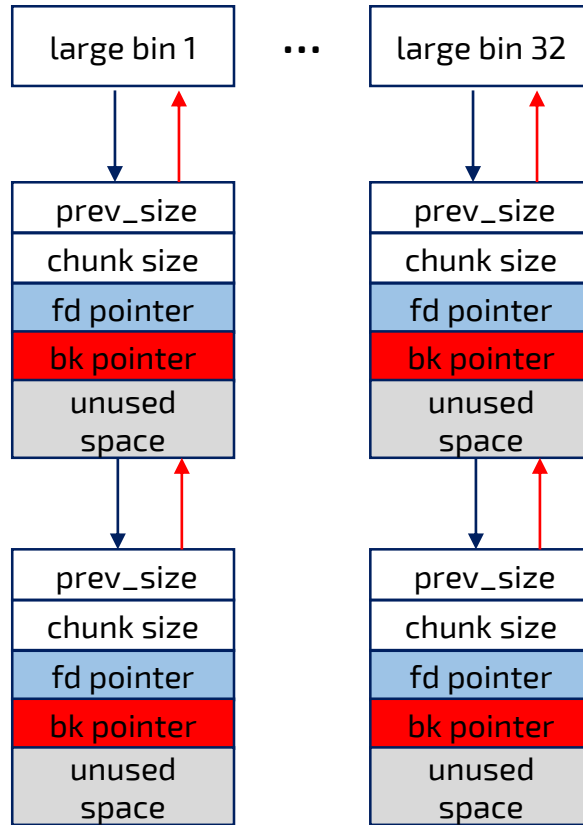


LARGE BINS


32x spacing: 64




16x spacing: 512



2x spacing: 256k




all else



UNSORTED BIN

- Rationale
 - » *often **freed** objects are clustered together*
- Strategy
 - » *merge them before placing them in the correct bin*
- Heap manager merges neighboring **deallocations**
 - » *and places them in the **unsorted bin***
- **malloc**
 - » *bins checked for good fit:*
 - » *if yes*
 - allocate it
 - » *otherwise*
 - place in correct small/large bin



FAST BINS

- Keep small, recently released chunks in a “fast” queue
 - » *sizes: 16, 24, 32, 40, 48, 56, 64, 72, 80, 88*
 - » *no neighbor merges*
 - P bit remains 0
- Periodically, heap mgr consolidates the heap
- Triggered by
 - » *a large **malloc** request*
 - » *freeing chunks over 64KiB*
 - » ***malloc_trim** / **mallopt** invocations*

TCACHE (PER-THREAD CACHE) BINS

- Problem

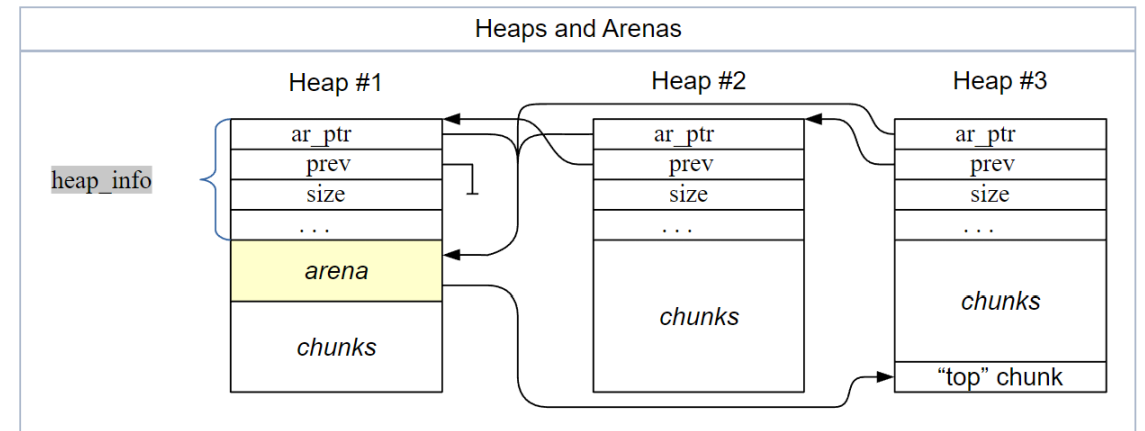
- » *lock contention in multi-threaded apps over shared resources*
- » *cost too high for heap allocations*

- Solution

- » *per-thread **arenas***
- » *each thread has 64 tcache bins*
 - up to 7 same-size chunks per bin
 - 12-516 / 24-1032 bytes

- Operation

- » *similar to fast bin*
- » *with slow-path default (acquire arena lock)*



THE FULL MALLOC

1. If **tcache** chunk available → return it
2. If request is very large → mmap it
3. Otherwise, acquire arena lock and
 1. *Try fast bin/small bin recycling*
 - If corresponding fast bin exists, try to find chunk there; pre-fill tcache opportunistically
 - Otherwise, if corresponding fast bin allocate from there
 2. *Resolve all deferred **fre**es*
 - “truly free” fastbin entries and consolidate to the unsorted
 - search unsorted for suitable; if found, stop. otherwise put entry in small/large/tcache
 3. *Default to basic recycling*
 - If the chunk size corresponds with a large bin, search the corresponding large bin now
 4. *Create new chunk from scratch*
 - request heap extension if not enough space
 5. *If all else fails, return NULL*

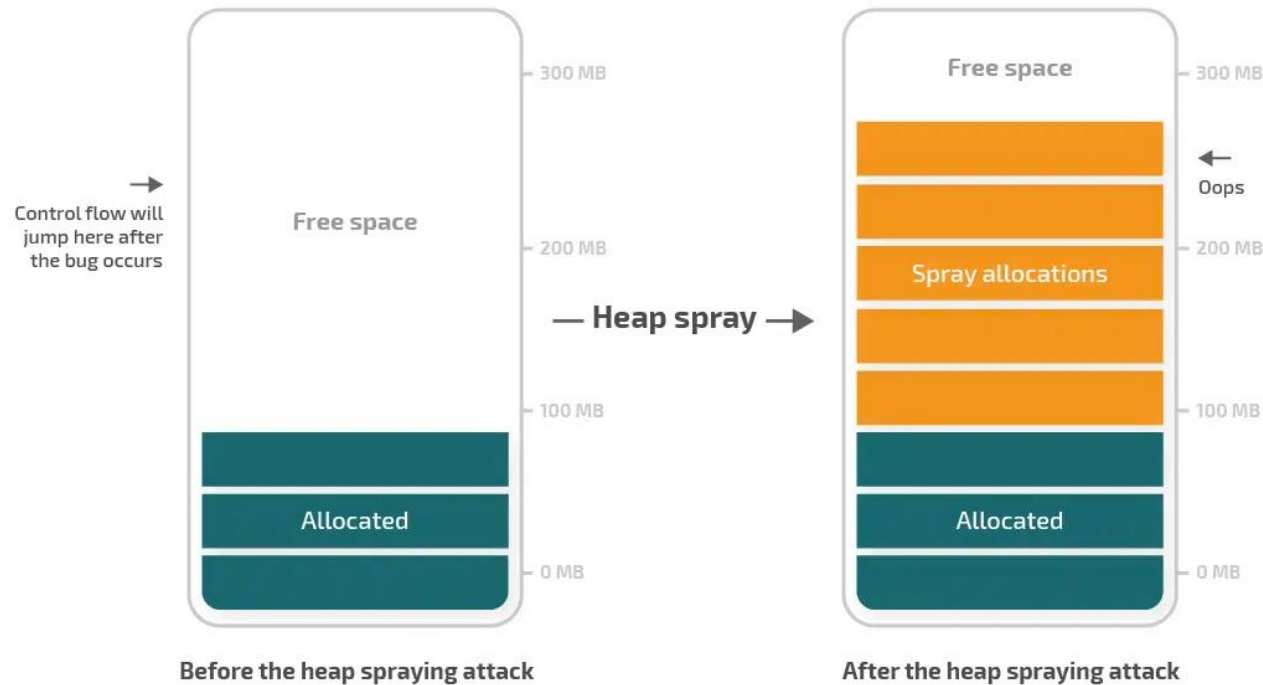
THE FULL FREE

1. If `ptr == NULL` → do nothing
2. Convert `ptr` to chunk
3. Perform sanity checks; abort on failure
4. If chunk fits in **tcache**, store it there
5. If chunk has **M** bit set, use **munmap**
6. Otherwise, obtain arena heap lock and
 1. *If chunk fits into a **fastbin** → place it there and return*
 2. *If chunk > 64K, consolidate **fastbins** and place merged chunks on the **unsorted** bin*
 3. *Merge the chunk backwards + forwards in **small**, **large** and **unsorted** bins*
 4. *If chunk at the top of the heap → merge into the **top***
 5. *Otherwise, store it in the **unsorted** bin*

REVISITING HEAP RULES / VIOLATIONS

- **Do not** read/write to a **malloc**-ed pointer after it has been **free**-d
 - » *Use after free vulnerability*
- **Do not** use or leak uninitialized info in a heap allocation
 - » *Information leaks & uninitialized data vulnerabilities*
- **Do not** read/write bytes after the end of an allocation
 - » *Heap overflow & read beyond bounds vulnerabilities*
- **Do not** pass a pointer that originated from **malloc** to **free** more than once
 - » *Double free vulnerability*
- **Do not** read/write bytes before the beginning of an allocation
 - » *Heap underflow vulnerability*
- **Do not** pass a point that did not originate from **malloc** to **free**
 - » *Invalid free vulnerability*
- **Do not** use a pointer returned by **malloc** before checking against **NULL**
 - » *Null-dereference bug & occasional arbitrary write vulnerability*

HEAP SPRAYING



```
nops = unescape('%u9090%u9090');  
s = shellcode.length + 50;  
  
while (nops.length < s)  
    nops += nops;  
f = nops.substring(0, s);  
block = nops.substring(0, nops.length - s);  
  
while (block.length + s < 0x40000)  
    block = block + block + f;  
  
memory = new Array();  
for (counter = 0; counter < 250; counter++)  
    memory[counter] = block + shellcode;  
  
ret = '';  
for (counter = 0; counter <= 1000; counter++)  
    ret += unescape("%0a%0a%0a%0a");
```

HANDS-ON: HOW2HEAP

USING HOW2HEAP EXAMPLES BY SHELLPHISH

`git clone https://github.com/shellphish/how2heap.git`

- Narrative

Attack	Target	Technique
First Fit	This is not an attack, it just demonstrates the nature of glibc's allocator	---
Double Free	Making <code>malloc</code> return an already allocated fastchunk	Disrupt the fastbin by freeing a chunk twice
Forging chunks	Making <code>malloc</code> return a nearly arbitrary pointer	Disrupting fastbin link structure
Unlink Exploit	Getting (nearly)arbitrary write access	Freeing a corrupted chunk and exploiting <code>unlink</code>
Shrinking Free Chunks	Making <code>malloc</code> return a chunk overlapping with an already allocated chunk	Corrupting a free chunk by decreasing its size
House of Spirit	Making <code>malloc</code> return a nearly arbitrary pointer	Forcing freeing of a crafted fake chunk
House of Lore	Making <code>malloc</code> return a nearly arbitrary pointer	Disrupting smallbin link structure
House of Force	Making <code>malloc</code> return a nearly arbitrary pointer	Overflowing into top chunk's header
House of Einherjar	Making <code>malloc</code> return a nearly arbitrary pointer	Overflowing a single byte into the next chunk