



CSCI 4621
INTRO CYBERSECURITY

OS Review: Processes & Address Spaces

Vassil Roussev

vassil@cs.uno.edu



Operating System Services

OS user
interface



GUI



CLI/shell

Applications



user

System call interface

kernel

Program
execution

Device I/O
operations


File
systems

Network
communications

Error
handling

Protection
& security

hardware



Core OS Services

≡ Program execution

- Load program code from HDD into RAM
- Run/execute the code
- End execution

≡ Device I/O

- Provide access to various devices
 - USB, ATA, PCIe, etc.

≡ File systems

- Provide (abstract) access to storage devices
 - HDD, CD, DVD, B-Ray, etc.

Core OS Services (2)

≡ Network & inter-process communications

- Allow exchange of data b/w processes and via the network

≡ Protection & security

- Provide means to manage and enforce security policies

≡ Error handling

- Detect and recover from hardware and software errors
- System debugging

Core OS Services (3)

≡ Accounting & resource management

- Typically split across other core services
- Keep track of resource usage & statistics
- Manage & enforce resource restrictions

UI Services

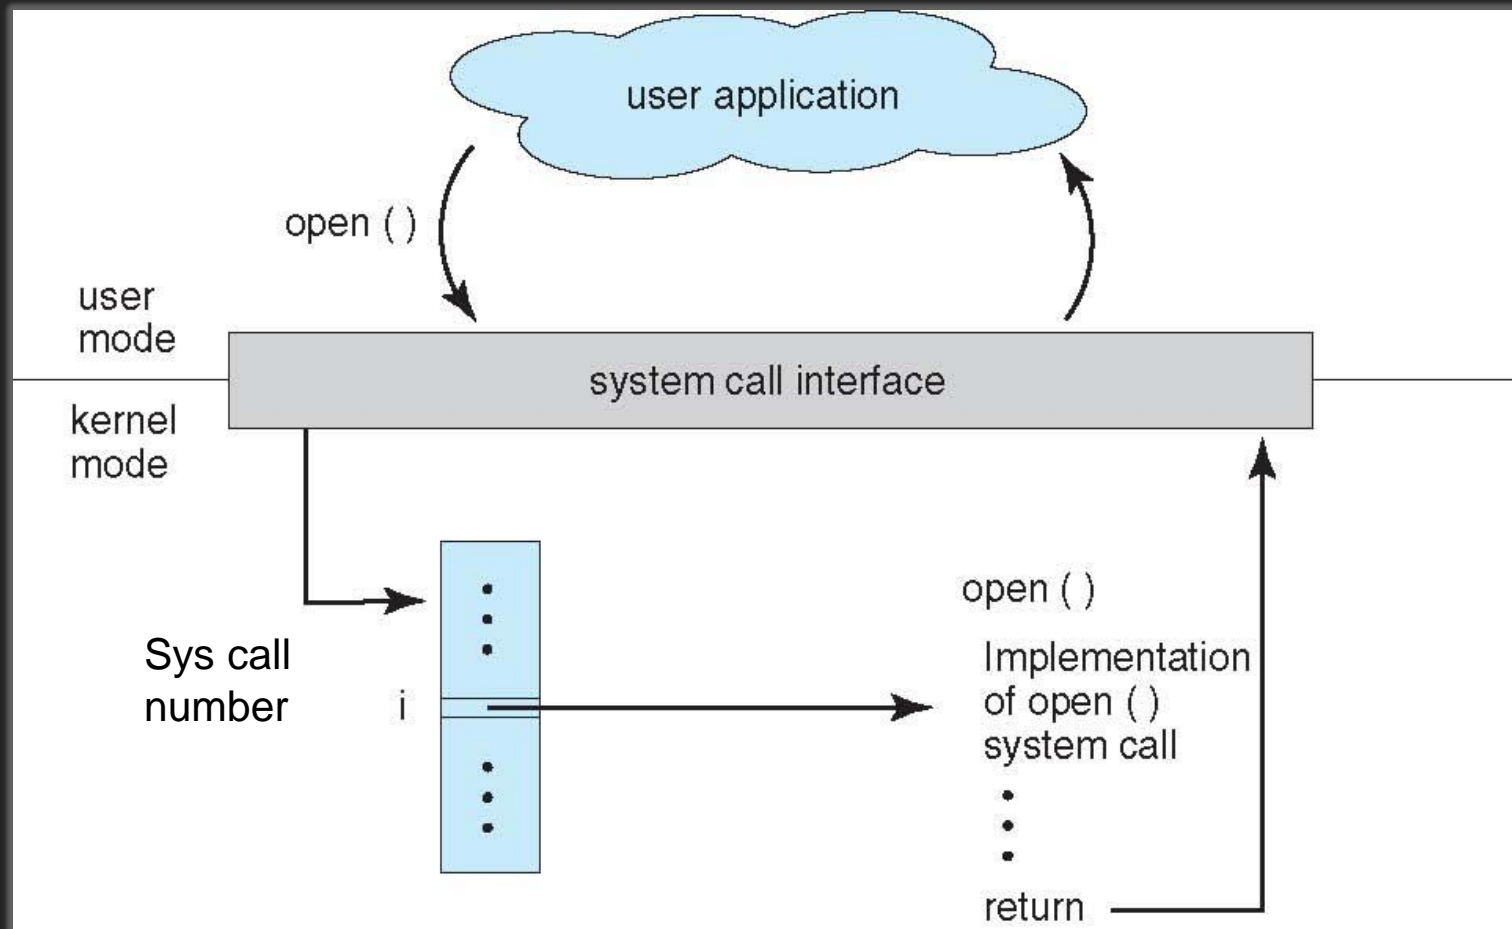
≡ These are (system) applications that exercise the system call interface

- They can be replaced with no ill effects

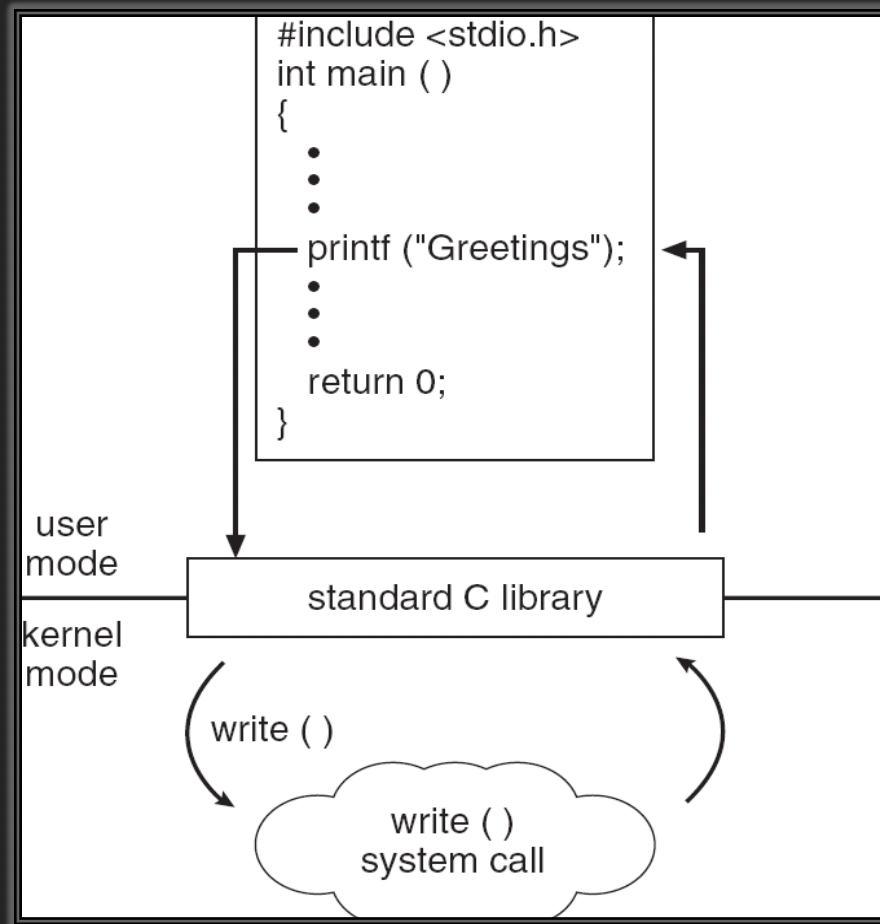
≡ UI

- GUI: graphical interface to
 - invoke/terminate programs,
 - manage user data,
 - manage access control, etc.
- CLI: command line interface
 - textual interface to access OS services
 - allows scripting/automation

API – System Call – OS Relationship



Sys call example: standard C library



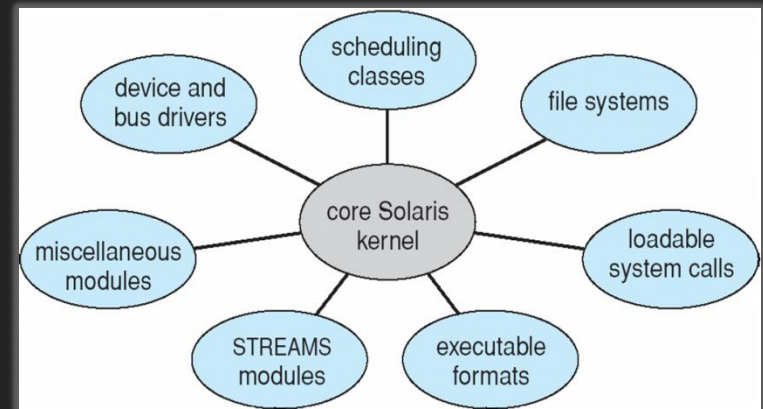
Sys call interfaces: Windows & Unix (POSIX)

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

Kernel Modules

≡ KM allow (to various extents)
modular extension of the
kernel

- Object-oriented approach
- Each core component is separate
- Each talks to the others over known interfaces
- Each is loadable as needed within the kernel




e.g.: **Solaris**

Virtual Machines

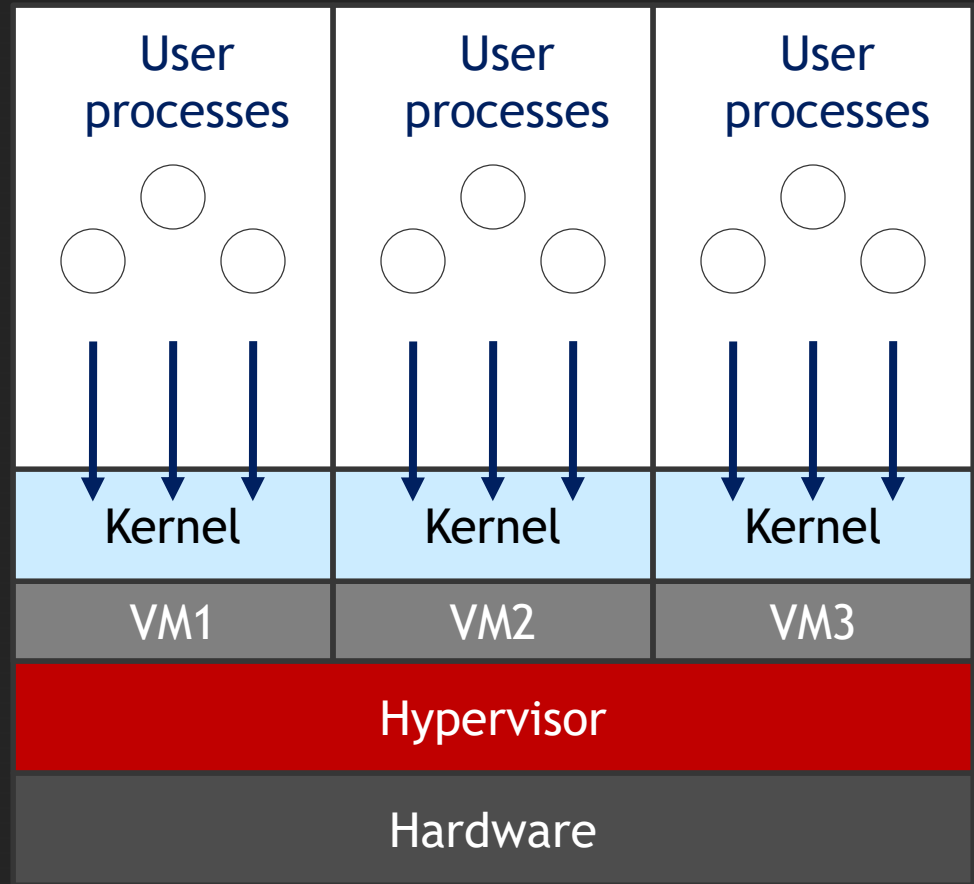
≡ A *virtual machine* provides an interface identical to the underlying bare hardware.

- The *hypervisor* creates the illusion that a process has its own CPU, RAM, I/O devices.
- Each guest provided with a (virtual) copy of underlying computer.
- Hardware support is needed for VMs to run efficiently

Virtual Machines

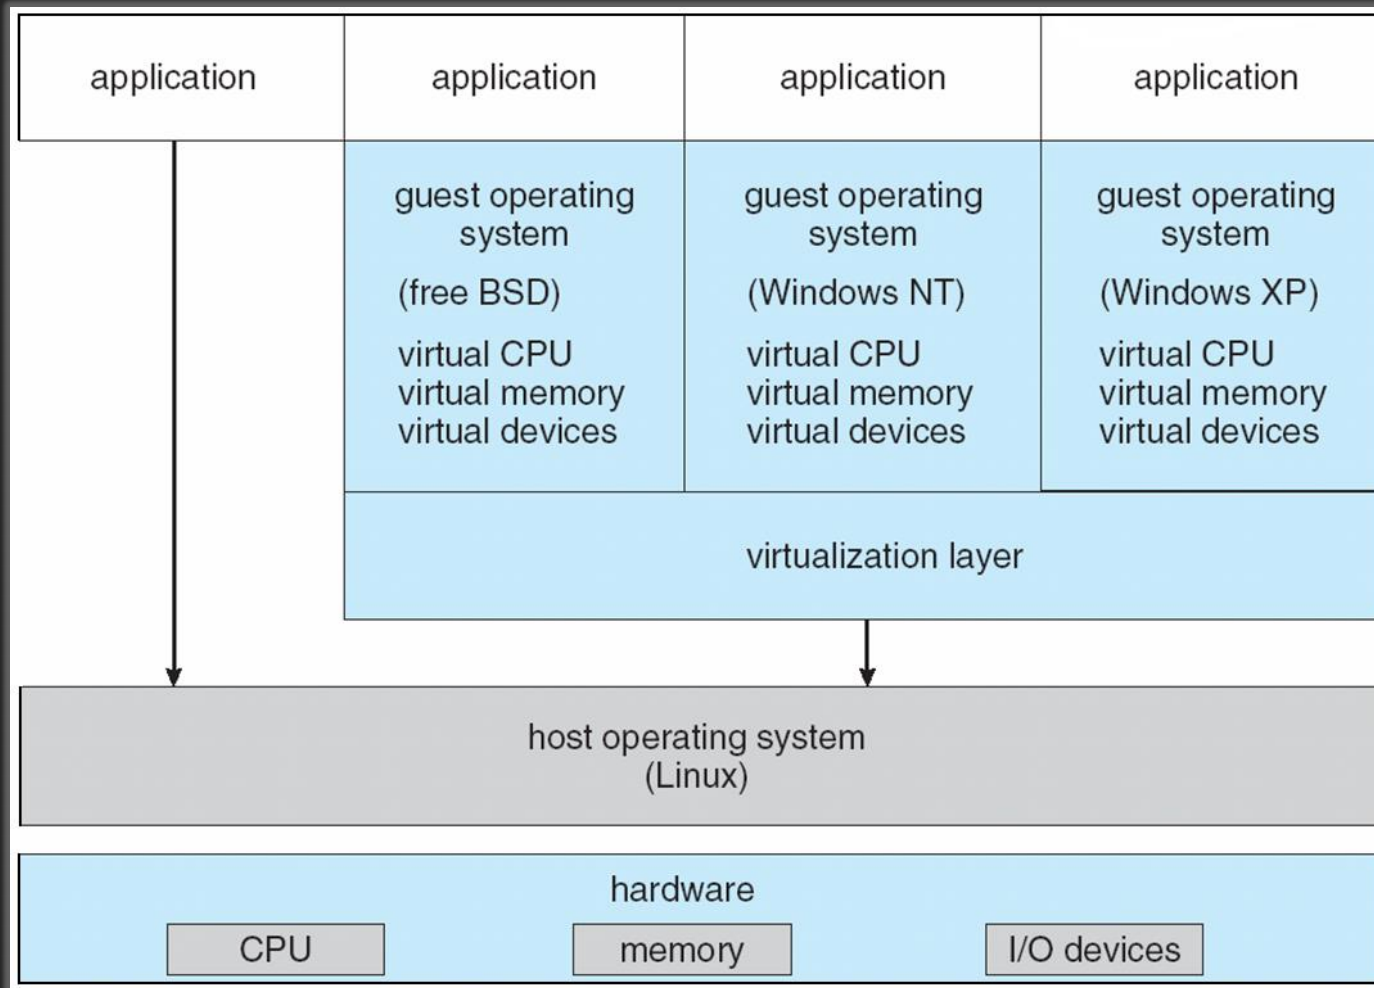


Non-virtualized machine:
Single OS instance is present



Virtualized machine:
Multiple OS instances may be present

Example VMWare Deployment



Linux: modular monolithic kernel

≡ All of kernel is one process

- all kernel components have access to everything

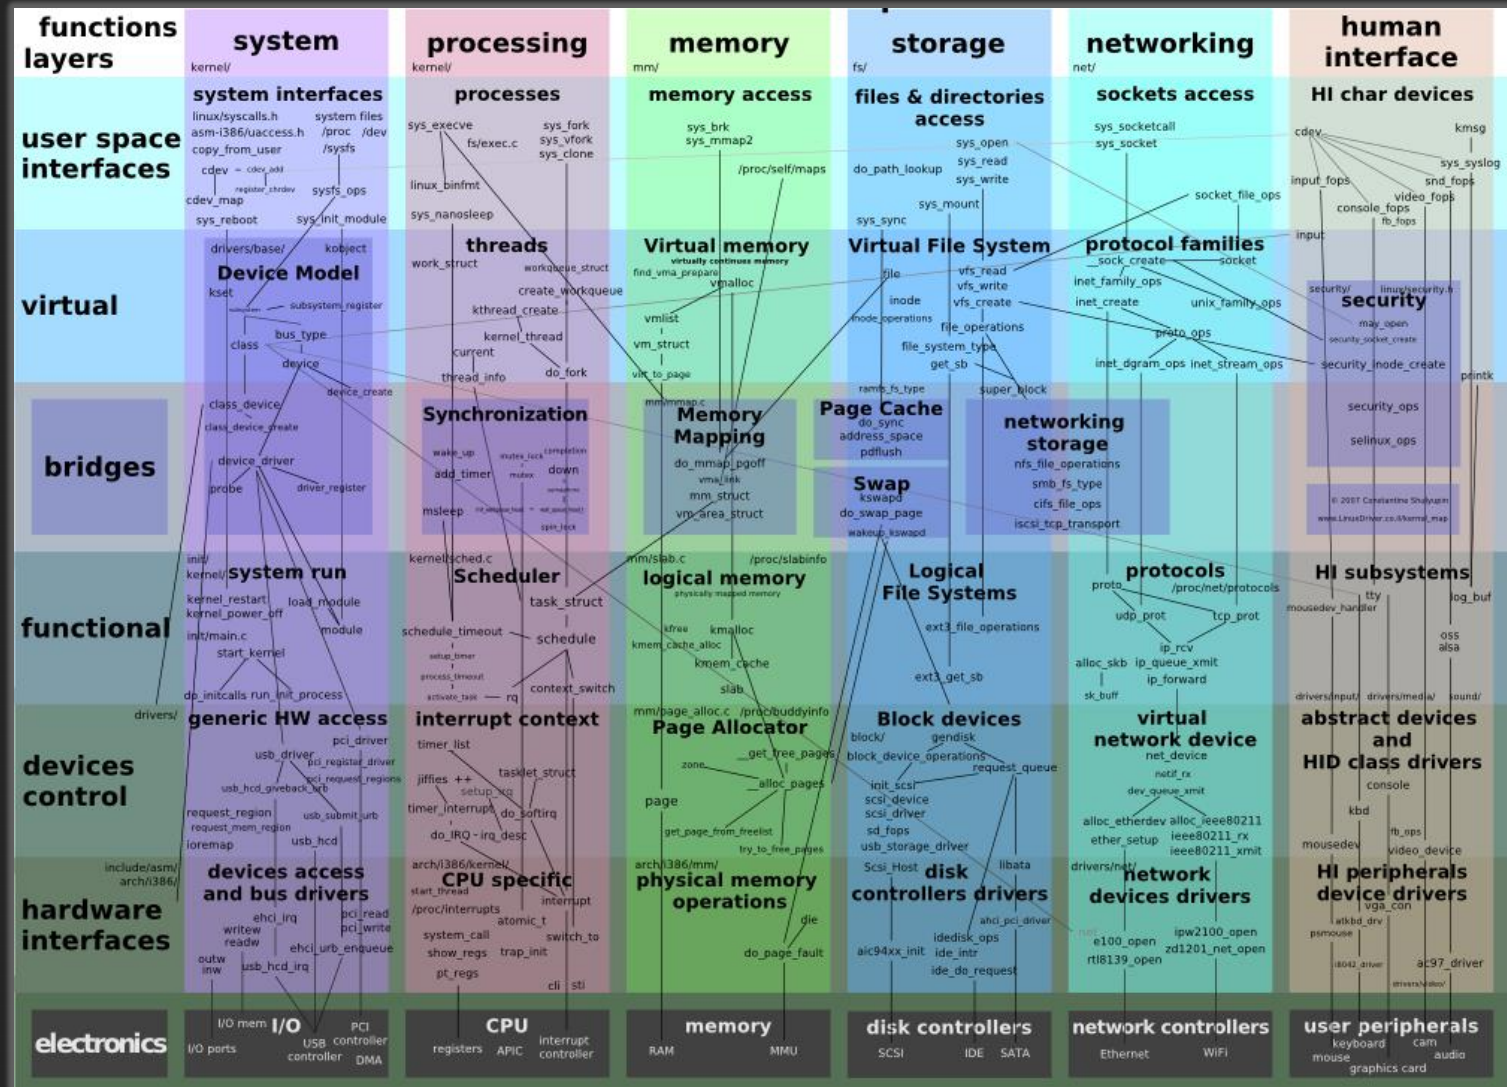
≡ Loadable kernel modules (LKM)

- autonomous blocks of code
- can be loaded/unloaded at run time
 - either explicitly, or as necessary

≡ Module stacking

- modules are arranged in a hierarchy
- modules may serve as libraries for other (client) modules

Linux kernel map



PROCESSES

Process definition

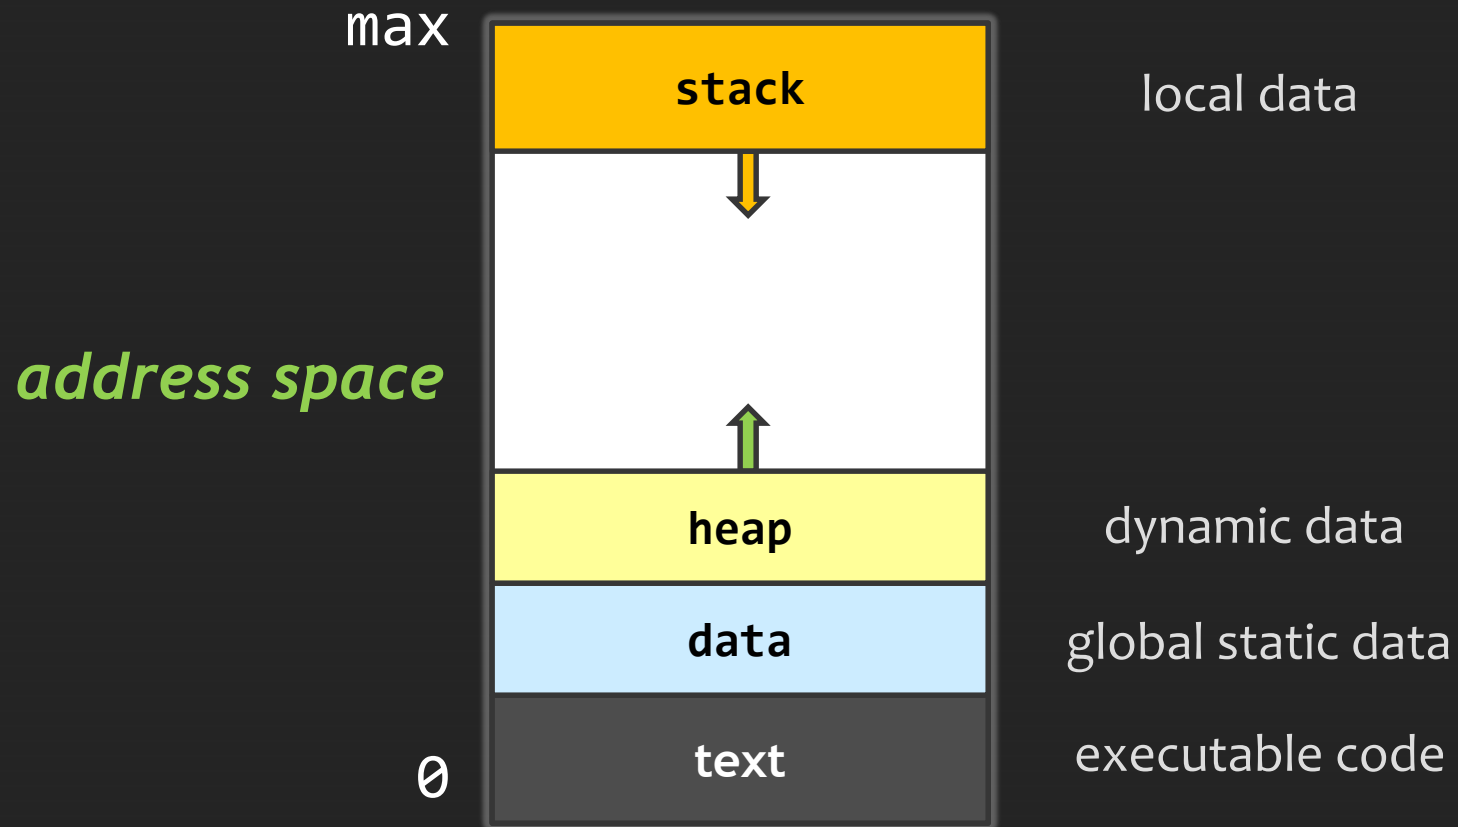
≡ A *program* is

- a sequence of instructions (a computation) that accomplishes a specific purpose.
- stored on persistent storage, loaded and executed when necessary.

≡ A *process/task* is a program in execution.

- each process gets a (dynamic) allocation of the system's CPU, RAM, and other resources.
- multiple instances of the same program can co-exist at the same time.
- by default, each process has its own address space that is protected from those of other processes.
- the number of processes executing in parallel is limited by available CPUs/cores.

Process in memory



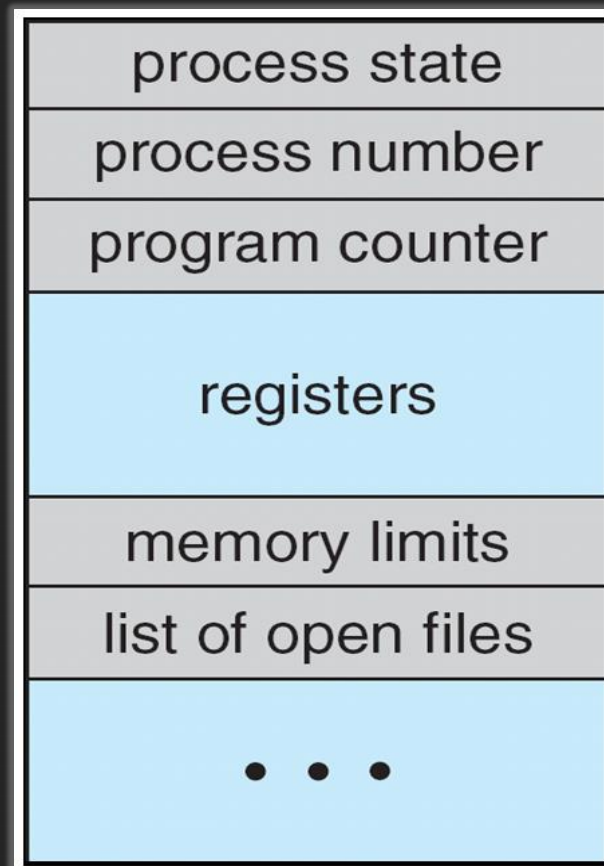
Process context

≡ Process context consists of all the information that describes the exact state of the process' computation.

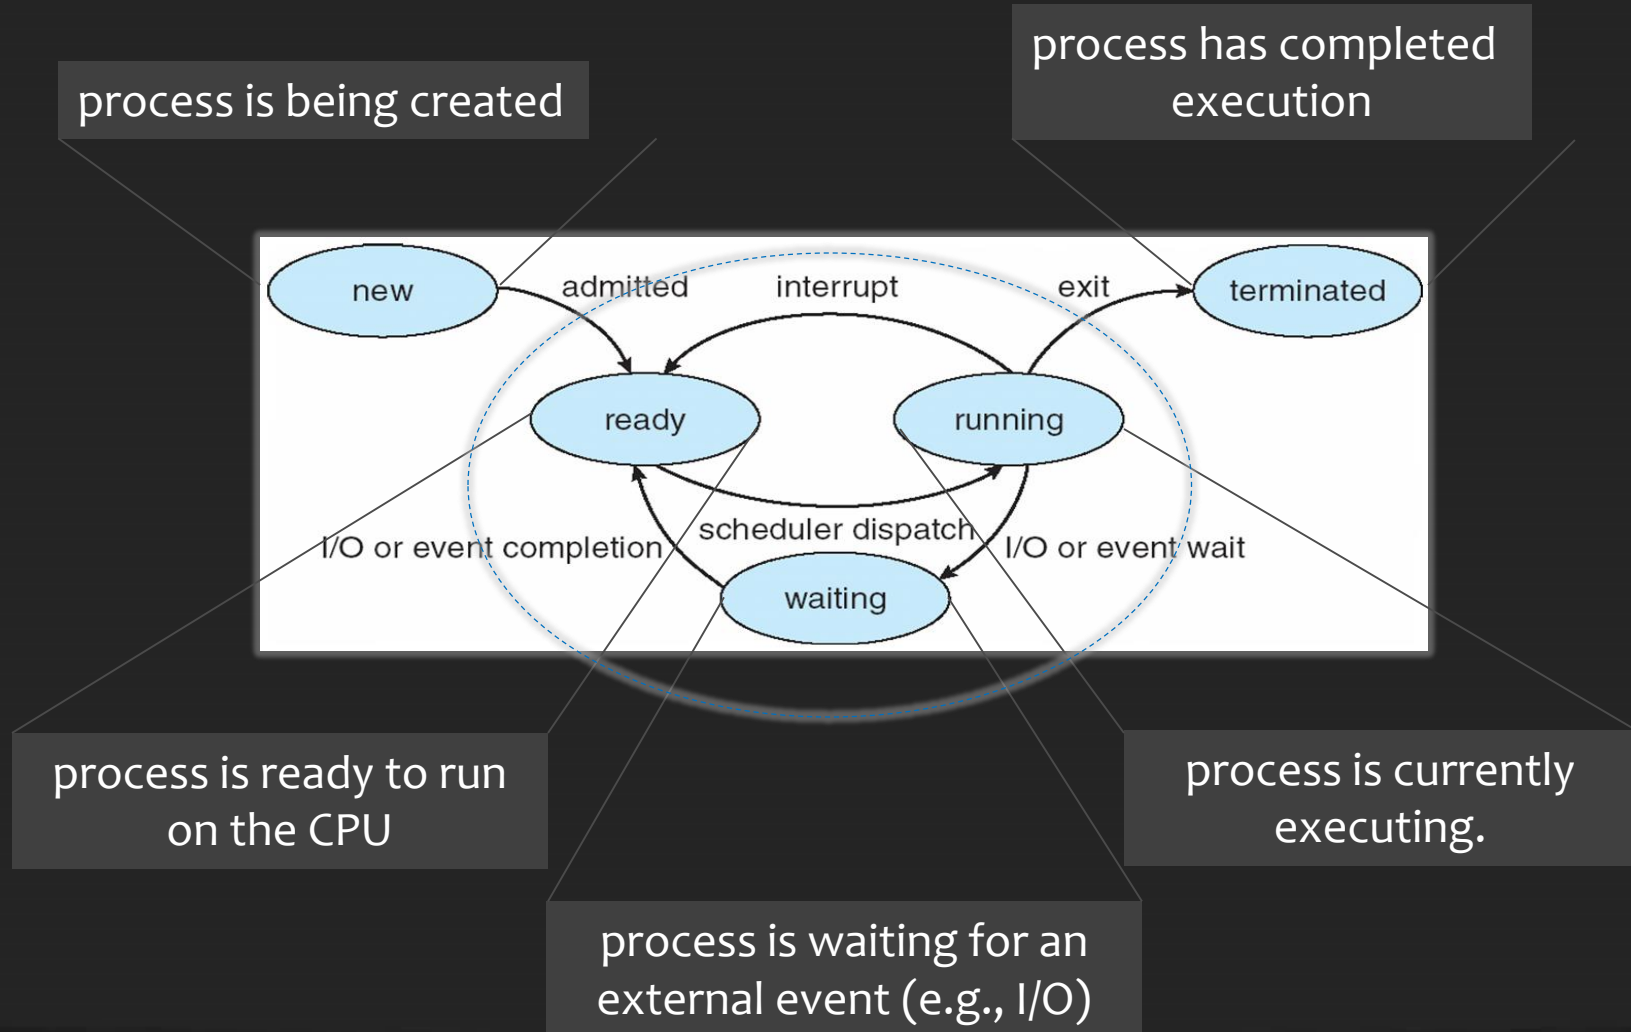
- the context consists of the contents of CPU registers and some scheduling info.
- if we take a snapshot of the process' state, we can freeze it and restart it later.

Process/task Control Block

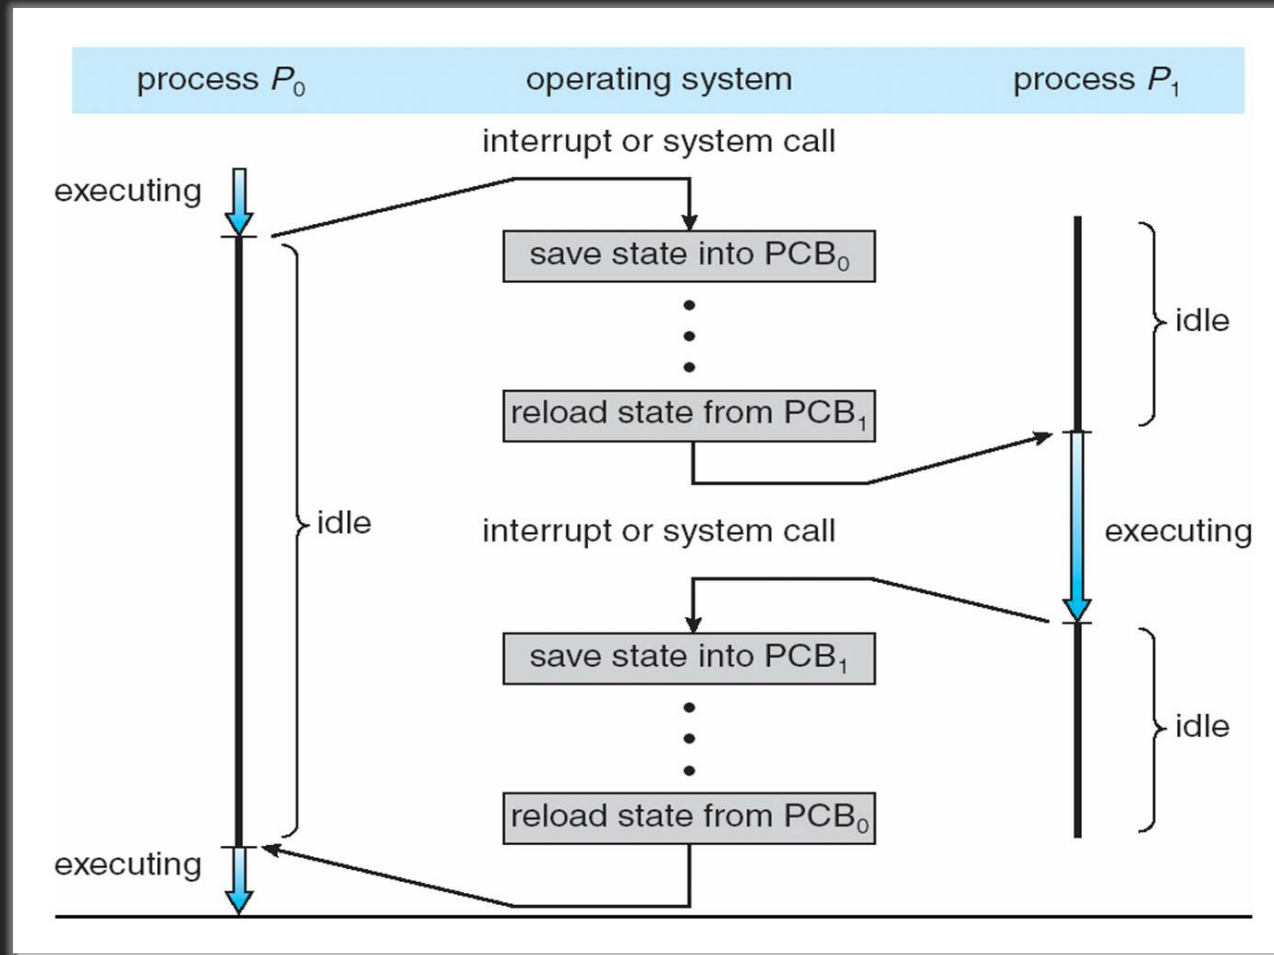
≡ Stores the process' context information:



Process states



Context switch



Context switch w/ cooperation

```
P0 () {  
    // some processing  
    ...  
    yield();  
    // more processing  
    ...  
}
```

```
yield () {  
    // find next process  
    // to be scheduled  
    ctxsw(old, new);  
}
```

```
P1 () {  
    // some processing  
    ...  
    yield();  
    // more processing  
    ...  
}
```

Minimal context switch: MIPS R4000

```
.globl ctxsw
.ent ctxsw
ctxsw:
    sw $2, 8($4)
    sw $3, 12($4)
    ...
    sw $14, 56($4)
    sw $16, 64($4)
    ...
    sw $25, 100($4)
    sw $29, 116($4)
    sw $30, 120($4)
    sw $31, 124($4)
    sw $31, 128($4)
    mflo $15
    sw $15, 132($4)
    mfhi $15
    sw $15, 136($4)

    s.d $f0, 140($4)
    ...
    s.d $f30, 260($4)
```

```
    lw $2, 8($5)
    lw $3, 12($5)
    lw $4, 16($5)
    lw $6, 24($5)
    ...
    lw $14, 56($5)
    lw $16, 64($5)
    ...
    lw $25, 100($5)
    lw $29, 116($5)
    lw $30, 120($5)
    lw $31, 124($5)
    lw $15, 132($5)
    mtlo $15
    lw $15, 136($5)
    mthi $15
    l.d $f0, 140($5)
    ...
    l.d $f30, 260($5)
    lw $15, 128($5)
    lw $5, 20($5)
    j $15
.end ctxsw
```


PCB in Linux

```
pid_t pid;           /* process identifier */
long state;          /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
...
```



CPU scheduling & dispatch

≡ Scheduling

- the algorithm by which CPU time is allocated to processes.

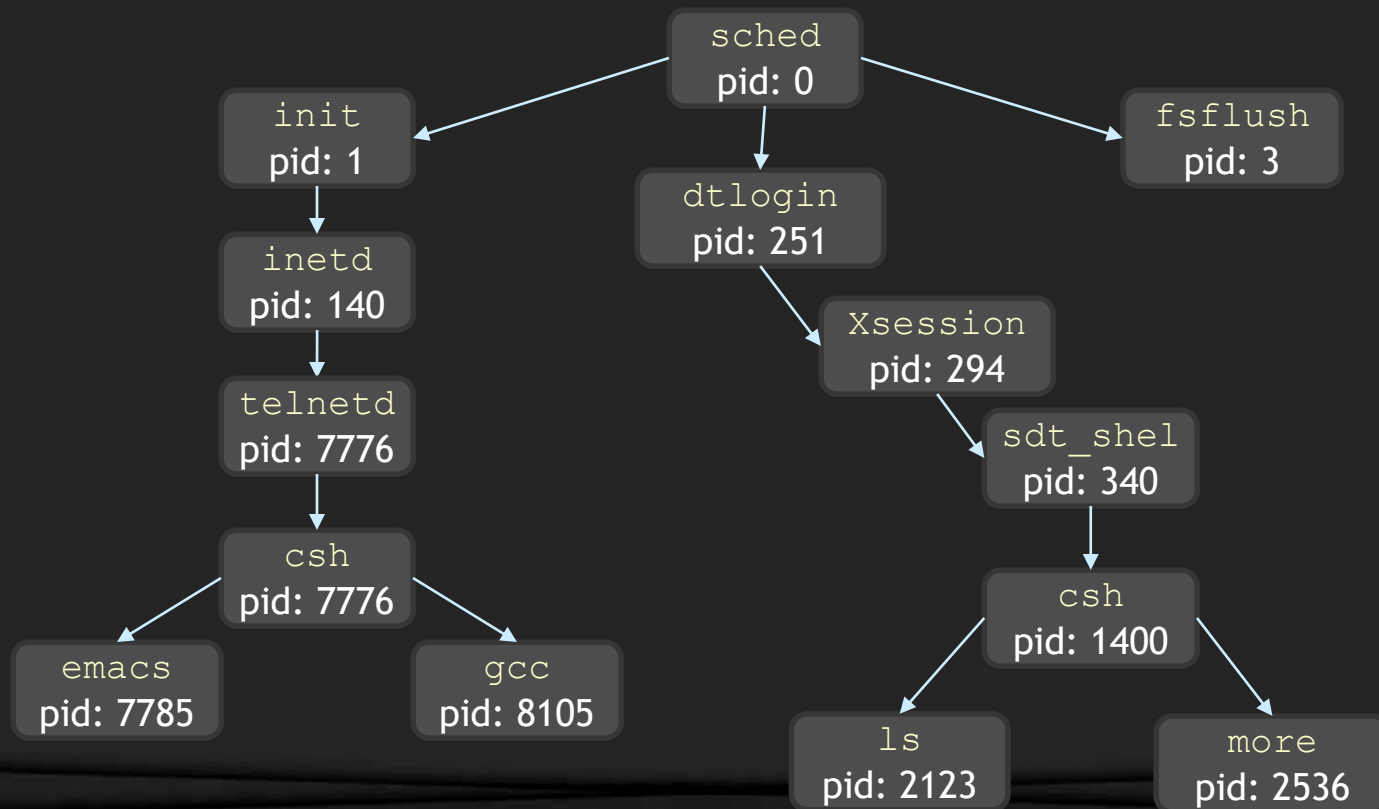
≡ Dispatch

- the procedure concerned with the immediate switch to a process selected for execution
 - context switch
 - switch to user mode
 - restart of the process


Process (pro)creation

≡ Hardware creates first process (pid 0)

≡ Remaining processes are *forked/spawned*; e.g.:




Forking/spawning a process



UNIX fork() example

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
```

```
int main() {
    pid_t pid;
    pid = fork();          /* fork another process */
    if (pid < 0) {         /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    } else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    } else {              /* parent process */
        wait (NULL);      /* wait for the child */
        printf ("Child Complete");
    }
    return 0;
}
```




Inter-Process Communication (IPC)

≡ By default, processes are isolated from each other; yet data exchange is often necessary for:

- resource sharing:
 - allow files, cached objects, etc. to be shared
- computational speedup
 - allow parallel computation
- modularity
 - separation of concerns
 - fault/security isolation

IPC models


Message Passing



Copy of data is passed along

- + more generic & secure
- less efficient (memory copy)

Shared Memory



Reference to the data is passed

- + more efficient
- requires trust & synchronization

Process synchronization & IPC

≡ **Synchronization** is any explicit order imposed on the execution of processes.

≡ *Blocking IPC* is a.k.a. *synchronous*:

- blocking **send()** stops sender execution until receipt by receiver
- blocking **receive()** stops receiver until message is available


≡ *Non-blocking IPC* is a.k.a. *asynchronous*:

- non-blocking **send()** returns immediately before result is known
- non-blocking **receive()** returns immediately either with waiting message, or **null**

≡ *Note:*

- synchronous/asynchronous is independently applicable to **send()/receive()**

Buffering and synchronization (local)



≡ Producer-consumer system

≡ Single shared buffer

- $|buffer| = 0 \rightarrow$ blocking send/receive, a.k.a. rendezvous
- $|buffer| = k < \infty \rightarrow$ up to k asynchronous send/recv
- $|buffer| = \infty \rightarrow$ completely asynchronous system—unlimited asynchronous send/recv

Buffering and synchronization (remote)



≡ Separate send/receive buffers

- $|buffer_1| + |buffer_2| = 0 \rightarrow$ blocking send/receive, a.k.a. rendezvous
- $|buffer_1| + |buffer_2| = k_1 + k_2 < \infty$
 - ➔ up to $k_1 + k_2$ asynchronous send
 - ➔ up to k_2 asynchronous recv

Local communication mechanisms

≡ Shared address space

- parent/child processes

≡ Shared segments


≡ Ordinary pipes

- unidirectional buffered communication
- requires parent/child relationship

≡ Named pipes

- bi-directional communications
- no parent/child relationship necessary

A socket is the basic **remote communication** abstraction provided by the OS to processes.




Network access is managed by the OS
→ OS provides a **socket** API to allow network communication

Socket programming using TCP service

Application Viewpoint

TCP provides reliable, in-order transfer of a stream of bytes between two end points (processes).

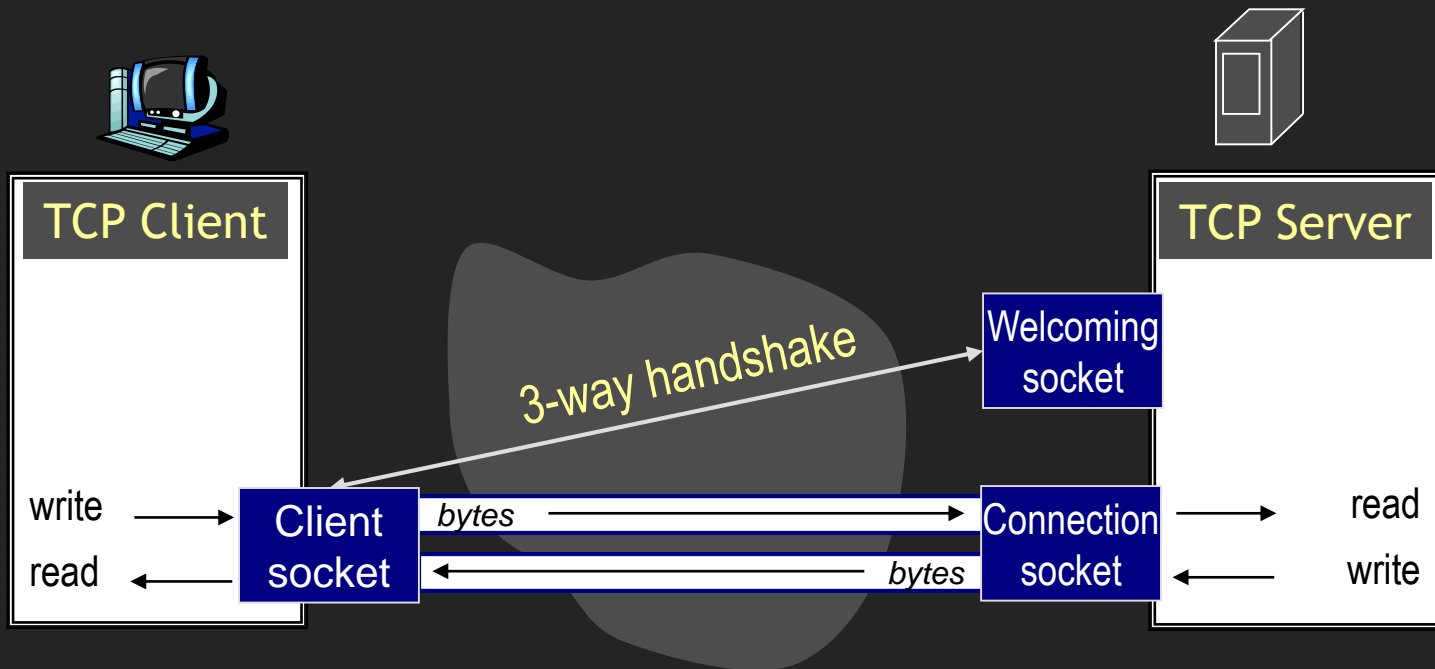


Same interface as other stream I/O
(files, pipes, etc.).

Pair of pipes abstraction

TCP communication is based on the client/server concept.

Server is started first & waits for connection requests.



Client initiates connection via 3-way handshake.

Once established, connection is completely symmetrical.

A port number is a 16-bit number that uniquely identifies a network process on a host.

Client must know port number to reach server.

Server *binds* to a desired port number (welcoming socket).

Most standard services have a default port:

File Transfer →

Secure Shell →

Simple Mail Transfer →

World Wide Web →

Print Service →

bind()

Port	PID
0000	---

...

0021	○
0022	○

0025	○
------	---

...

0080	○
------	---

...

0631	○
------	---

...

ftpd

sshd

smtpd

httpd

cupsd

Client sockets use any available port >1023.

Network addressing for sockets

A ⇔ B socket connection: <IP-hostA, portA, IP-hostB, portB>

